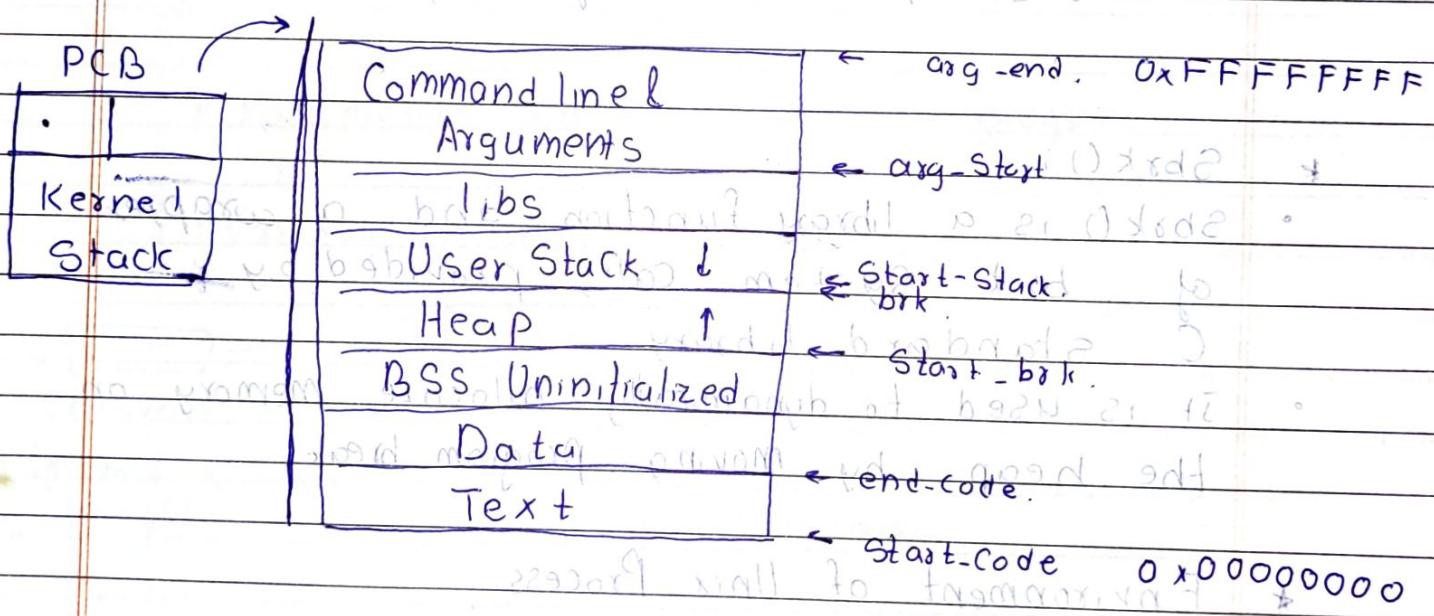


* mm_struct and memory layout of a process.

- In Linux, each process has an associated `mm_struct` which contains all the information about the process's virtual memory layout.
- It is defined in the kernel source (`include/linux/mm_types.h`) and is crucial for managing memory, performing context switches and handling system calls like `brk()`, `mmap()` etc.



`struct mm_struct {`

- `unsigned long Start-code, end_code;`
- `unsigned long Start-brk, brk;`
- `unsigned long Start-stack, start-mmap,`
- `unsigned long arg_start, arg_end`
- `unsigned long env_start, env_end.`

3.

System Call Used to terminate process.

-exit()

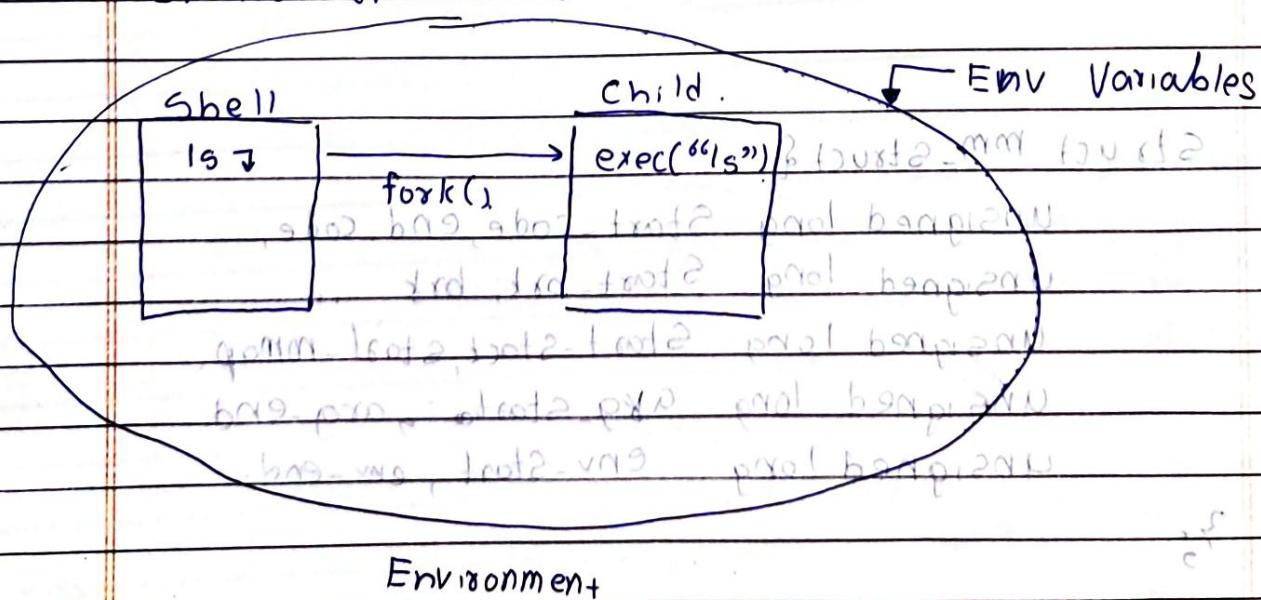
Page:
Date:

Start_Code // Start of Code Segment
end_code // End of Code Segment.
Start_brk // Start of heap
end_brk // Current end of heap
Start_stack // Start of User Stack
arg_start // Arguments
arg_end // End of Arguments
env_start // Environment
env_end // End of Environment
start_mmap // Base of mmap area.

* Sbrk():

- Sbrk() is a library function and a wrapper of brk system call, provided by C standard library.
- It is used to dynamically allocate memory on the heap by moving program break.

* Environment of Unix Process.



Shell (bash)

call creates new stream of begin at T

→ fork() → Child Process

→ dup → exec("ls")

→ New Process Image:

- argv = ["ls"]

- envp = inherited from

parent (Shell)

child has its own copy of memory layout etc

→ wait()

* Parameters Required to run program.

1. Variable known to the shell

global variable	PATH	root to execute
environ	HOME	home dir.
	USER	user name
	SHELL	/bin/bash

These are environment variables that are known to the shell

2. getenv()

It is a function in C used to retrieve the value of an env variable.

3. putenv()

The putenv() function is used to set or modify an env variable in the current process.

Who creates → init()
Shell.

first process
Date _____
Page _____

4. export Command.

It is used to mark env variables for

export ~~229089~~ child process.

ex. ~~229089~~ export ~~229089~~ MY_VAR=42

will export my program.

Ex. ~~229089~~ export

MY_VAR=42

will have

Set Command :

Displaying Shell variables and functions

* Signals:

- Signals are software interrupts used for asynchronous notifications.
- They inform processes of the occurrence of asynchronous events.
- Signals are a form of inter-process communication.
- Processes may send each other signals with a kill system call, or the kernel may send signals internally.

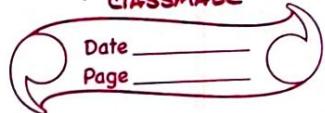
* Handling Signals.

There are three cases for handling signals: the process exits on receipt of the signal, it ignores the signal, or it executes a particular (user) function on receipt of the signal.

The kill() fun is used to send signals to processes.

CLASSMATE

int kill(pid_t pid, int sig);



Date _____

Page _____

- The default action is to call exit() in kernel mode, but a process can specify special action to take on receipt of certain signals with the signal system call.

Syntax:

```
oldfunction = signal(signum, function);
```

Where signum is the signal number the process is specifying the action for, function is the address of the user function the process wants to invoke on receipt of the signal.

oldfunction was the value of function in the most recent call to signal for signum.

When handling a signal the kernel determines the signal type and turns off the appropriate signal bit in the process table entry, set when the process received the signal.

```
int main() {
```

```
    signal(SIGINT, myfn);
```

```
    for (i=0; i<5000; i++) {
```

```
        printf("I am in loop %d\n", i);
```

```
        sleep(2);
```

```
    void myfn() {
```

```
        printf("Control-C was pressed");
```

```
}
```

Output: I am in loop after that adt loop will be end when user enter "exit" as stat of menu then Adt will be terminated after that Ctr+C was pressed

`(ctrl+c) → default action is to terminate`
`But the program doesn't terminate on`
`(ctrl+c) because we catch SIGINT`

How to find out Signals available in System

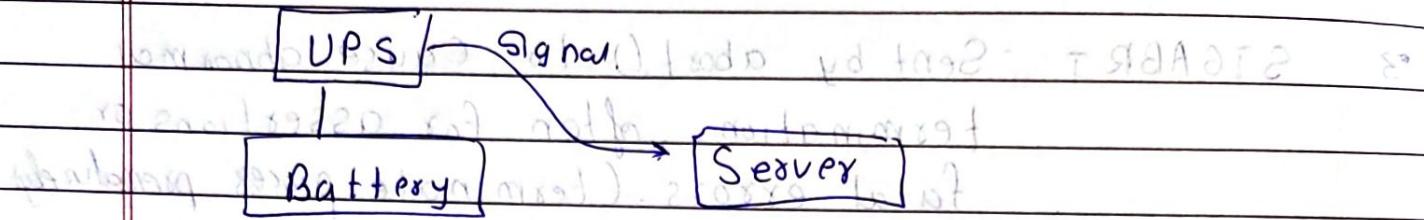
Killed and eaten on airband radio

return 0 in os return control to whom → to OS more specifically to parent process ex Shell like bash if you run program on terminal on fo

Standard & Common Signals

1. SIGALRM - Sent when a timer set by alarm() expires.
Used to implement time outs
 2. SIGQUIT - Sent when user presses Ctrl+
^] (Ctrl+backslash). and also
dumps core (for debugging).

- 3 SIGABRT : Sent by abort() to cause abnormal termination, often for assertions or fatal errors. (terminate process prematurely)
- 4 SIGURG : Indicates urgent data is available on Socket. Used in networking (TCP/IP)
- 5 SIGPIPE : Sent when a process writes to a pipe with no reader. Prevents broken pipe issues. (93BSD)
- 6 SIGCHLD : Sent to a parent when a child process terminates or stops. (Srv) Used in wait() handling.
- 7 SIGBUS : Indicates bus error, like unaligned memory access or invalid physical address.
- 8 SIGSYS : Sent when a process calls an invalid or restricted syscall.
- 9 SIGCONT : Continue a stopped process.
- 10 SIGEMT : Emulation Trap
- 11 SIGFPE : Floating Point Exception
- 12 SIGHUP : Hangup on terminal disconnected.
- 13 SIGINT : Interrupt from keyboard (Ctrl+C)
- 14 SIGIO : Some events occur in IO.
- 15 SIGKILL : kill the process immediately - Cannot be caught or ignored.
- 16 SIGPOLL : From pollable device (Pollable event)
- 17 SIGPROF : Profiling timer expired Set by Setitme.
- 18 SIGPWR : Power failure or battery low.



- 18. SIGQUIT - Abnormal termination or break -
- 19. SIGSEGV - Segment Violated (Segmentation fault) -
- 20. SIGSTOP - Stop the process immediately -
Can't be caught or ignored.
- 21. SIGTERM - Termination request.
- 22. SIGTRAP - Trap - Used by debuggers.
- 23. SIGSTP - Terminal Stop (Ctrl+Z)
- 24. SIGTTIN - Background process tried to read from terminal.
- 25. SIGTTOUT - Background process tried to write to terminal.
- 26. SIGUSR1, 2 - User-defined Signals.
USR 2 ?
- 27. SIGVTALRM - Virtual timer expired.
- 28. SIGWINCH - Window Size changed.
- 29. SIGXCPU - Sent when process exceeds its CPU time limit.
- 30. SIGXFSZ - Sent when process tries to write more data to a file than its file size limit.

```
#include <stdio.h>
#include <signal.h>
int main() {
    printf("One\n");
    signal(SIGALRM, alarm_catch); // define handler and
    alarm(10); // Set the timer
    pause(); // Pause program until a Signal is received
    printf("Two\n");
}
alarm_catch() {
    return;
}
```

Output: One

- We can implement functionality like `Sleep()` using `SIGALRM` `alarm()`.

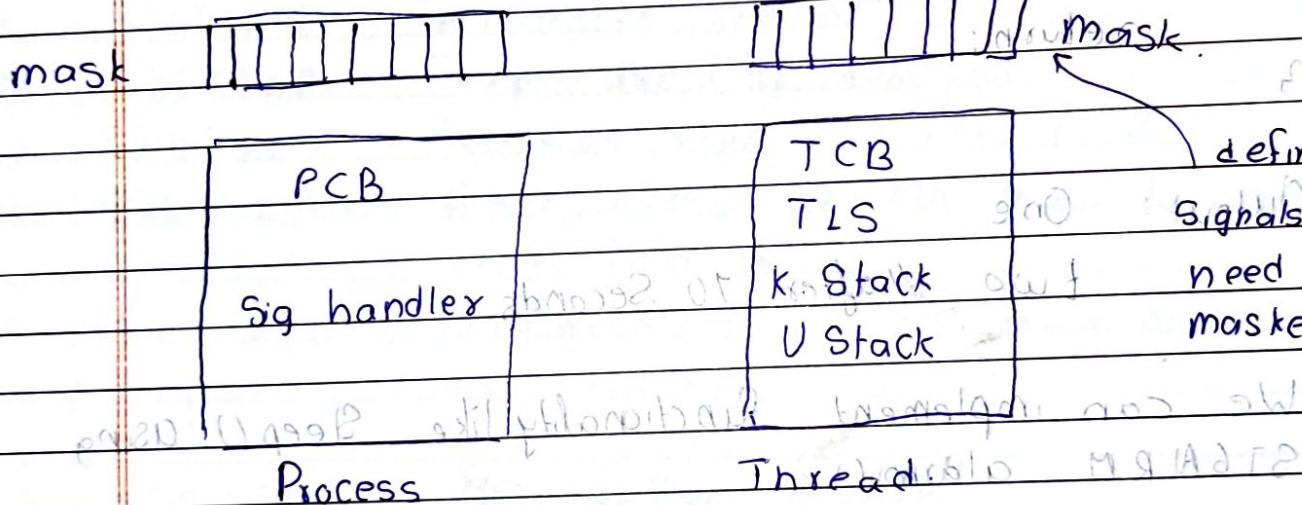
```
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
void segfault(int dummy) {
    printf("Segment Fault Occurred!\n");
    exit(1);
}
int main() {
    int * p = 0;
    signal(SIGSEGV, Segfault); // Signal handler
    *p = 1;
    return 0;
}
```

Output SegFault Occurred.

P is pointing to address 0. (Nullptr)

We are trying to store
17 at memory location of P (which
is address 0, not allowed).

The address does not belong to this process
∴ SegFault occurs.



Signal → Shared by both process and thread.

```
#include <signal.h>
```

```
int Sigaddset (sigset_t *Set, int signo);
```

```
Sigdelset (
```

```
Sigemptyset (
```

```
Sigfillset (
```

```
Sigismember (const sigset_t *Set, int no);
```

```
pthread_sigmask (
```

```
int Sigprocmask (int how, *sigset_t *newmask,
```

```
sigset_t *oldmask);
```

- * **Signal Set**: It is a data structure used to represent a collection of signals.
- 1. **Sigaddset**: Adds a specific signal to the set.
- 2. **Sigemptyset**: Initializes set to an empty set.
- 3. **Sigdelset**: Removes a signal from the set.
- 4. **Sigismember**: Tests if a signum is in set.
- 5. **Sigfillset**: Fills set with all signals.
- **pthread_sigmask**: Changes signal mask of the calling thread.

~~• **Sigprocmask**: Similar to **Sigmask** but for single threaded programs.~~

- **Sigaction()**: It is used to define a custom action that a process should take when a specific signal is received.

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#include <signal.h>
```

```
#include <unistd.h>
```

```
Static void Sigquit(int signo){}
```

```
main(){Caught SIGQUIT",
```

```
signal(SIGQUIT, SIG_DFL);
```

```
// restore default behavior for Sigquit.
```

```
// Next time it's received, the program will terminate
```

```

int main()
{
    sigset(SIGQUIT, sig_quit); // Set handler
    signal(SIGQUIT, sig_quit); // Set handler
    sigemptyset(&newmask); // Initialize empty set
    sigaddset(&newmask, SIGQUIT); // Add to set
    sigpromask(SIG_BLOCK, &newmask, &oldmask); // Block SIGQUIT
    printf("SIGQUIT blocked for 10 seconds.\n");
    sleep(10);
}

```

```

    {
        if (sigismember(&pendmask, SIGQUIT))
            printf("SIGQUIT is pending\n");
    }
}

```

```

    sigpromask(SIG_SETMASK, &oldmask, NULL); // Unblock SIGQUIT
    printf("SIGQUIT is unblocked\n");
    sleep(10);
    exit(0);
}

```

- We register a custom signal handler for SIGQUIT.
- We prepare a signal set containing only SIGQUIT.
- We block SIGQUIT for 10 seconds. If we press Ctrl+H now, the signal is not delivered, but it is marked as pending.
- Then we check for any pending signal.
- If we find SIGQUIT in pending we print SIGQUIT is pending.

- Then we unblock the signals. If the signal was pending, it's now delivered, and the handler `bsig_quit()` is executed.
- Inside `ss` the handler, it resets `SIGQUIT` back to default behavior (ie, the next time it's received, the program will terminate).

Output : SIGQUIT blocked for 10 seconds by pressing `Ctrl+D` now.

^D Pressed by user.

Caught SIGQUIT keys //After 10 sec.

SIGQUIT unblocked

//default behaviour

end of programmatic loop.

* Process Tracing

- Process tracing lets one process (like a debugger) observe and control the execution of another process.
- Typical things a tracing process can do:
 - Inspect or change registers and memory
 - Intercept system calls made by IT
 - Stop/resume the traced process
 - Inject breakpoints
- A debugger process, such as Sdb, spawns a process to be traced and controls its execution with the `ptrace` system call.

- A process (tracer) can observe and control another process (tracee).
- It is used for and can do the following:
 - Debugging.
 - Security auditing.
 - Sandboxing.
 - Reverse engineering.

* Working

1. Target process calls ptrace(PTRACE_TRACE_ME).
 2. Parent process calls wait() → waits for child to stop at next signal or exec()
 3. Parent now can use ptrace() system call to
 - read/write into child's memory (PEEK DATA, POKEDATA)
 - read/write registers (GETREGS, SETREGS)
 - Stop/resume execution (SINGLE STEP, CONT)
 - Intercept System call.
 4. Each time the tracee hits a syscall, signal, or trap, the tracer gets notified.
- * What is ptrace():
 - Ptrace is a System call that allows one process (usually a debugger) to observe and control execution of another process.
 - It is used by debuggers like gdb.
 - ptrace stands for process trace.

```

#include <stdio.h>
#include <sys/types.h>
#include <sys/user.h>
#include <sys/ptrace.h>

int main() {
    int pid = fork();
    if (pid == 0) { // Child process
        ptrace(PTRACE_TRACE_ME, 0, 0, 0);
        execve("/bin/ls", {"ls"}, 0);
    } else { // Parent process
        while (1) {
            int status;
            struct user_regs_struct regs;
            if (waitpid(pid, &status, 0)) {
                if (WIFEXITED(status))
                    printf("Child exited.\n");
                else if (WIFSTOPPED(status))
                    printf("Child has been stopped.\n");
                if (ptrace(PTRACE_GETREGS, pid, 0, &regs)) {
                    printf("eip=%lx, esp=%lx\n", regs.eip, regs.esp);
                    sleep(1);
                    ptrace(PTRACE_SINGLESTEP, pid, 0, 0);
                }
            }
        }
    }
}

Here, the child process calls PTRACE_ME
and executes /bin/ls where child is
stopped at every sys call and notified to parent.

```

Parent process waits for child using waitpid.

- Then we trace out execution of its code for
- We print different contents of registers. (extended instruction ptr, Extended Stack pointer)
- When child runs an instruction it stops and send signal to parent SIGTRAP.
- Every PTRACE_SINGLESTEP that generates a SIGTRAP.
- Ptrace Single Step tells child to execute one instruction and stop and send signal to parent.
- Child process is the process we need to debug.

* Syntax:

```
long ptrace(enum __ptrace_request request, pid_t pid,
           void *addr, void *data);
```

request: The operation to perform.

pid: The process Id of the target process

addr: Memory address, depending on the req

data: Data to write or a pointer to

receive data, rep depending on the request

• PTRACE TRACE ME:

Called by a child process to indicate that it wants to be traced by its parent.

Then parent can then use another ptrace

request to attach. Usually followed by exec() in child

• PTRACE ATTACH:

Attach to an already running process (target process is stopped). Similar to

what debugger's do when they attach to a PID.

Once you are done you can call CONT or DETACH to let it continue.

• PTRACE DETACH

Detach from a process that was previously attached. (Can optionally continue the process)

• PTRACE CONT:

Continue execution of a stopped process

Optionally, a signal can be sent to the process.

• PTRACE KILL

Sent SIGKILL to the traced process.

It kills the traced process (terminates).

• PTRACE SINGLESTEP

Resumes execution for one instruction and then stops (like a breakpoint step).

* Memory Access.

- **PTRACE_PEEKTEXT** ~~based at offset of address~~
Read a word from the code Segment (instruction Space) (~~stack, heap, global Data~~)
Used when you need to read machine instructions.

* PTRACE_PEEKDATA

- Read a word (usually 4 or 8 bytes) from the target's memory at offset (stack, heap, global variables etc.)
Used when you want to inspect variables, buffers, or runtime memory.
- **PTRACE_POKETEXT** ~~addr in tel of HATED~~
Write a word to the code Segment

- **PTRACE_POKEDATA** ~~addr of data~~
Write a word into target's memory (data Segment)
Used to modify variables etc.

- **PTRACE_PEEKUSER** ~~Implied in Ptrace~~
Read a word from the user Area (Registers) etc. (debug info.)

- **PTRACE_POKERUSER** ~~addr of data~~
Write a word to the user area

* Advanced.

- **PTRACE_SRESCALL**: Return of NOV 2022
 - Stop the trace at entry and exit of system calls.

Used to trace Systemcalls.

Used to intercept every system call - both when a call is entered and when it exits.

* Registers and CPU State.

• PTRACE_GETREGS

Get all general purpose registers (RAX, RSP, RIP).
Fetch all reg state

• PTRACE_SETREGS

Set general purpose registers.

Used to modify the CPU register value of tracee.

• PTRACE_GETFPREGS

Get floating point registers.

Allows you to read floating point registers of tracee.

• PTRACE_GETFPXREGS

Get extended floating point registers of tracee.

ex. SIMD, SIMD2, SIMD3, SIMD4

There are 16 general purpose registers.

Page: 6
Date: 1

- PTRACE_SETFREGS.
 - Set floating point registers.
 - Allows you to modify floating point registers of traced process to affect soft float mode.
- PTRACE_SETFPXREGS.
 - Set Extended floating point registers.
 - It is used to modify the extended floating point registers of a traced process.

e.g. SSE, AVX.

* What is User area? User area typically refers to the portion of memory allocated to a user process in OS.

* Typical ELF Relocatable Object file.

An ELF (Executable and Linkable Format)

relocatable object file consists of several

important sections and headers that

define how the file is organized.

An ELF file is a common standard file format used for executables, object code, shared libraries, and core dumps.

in Unix-like operating systems.

It is created by compiler (.o files), assembler, linker, debugger etc.

elf header	ELF file header	It is program header table for executable (optional).
sections	.text .data .bss .reltext .reldata .debug (local variables in this file)	read/write memory Segment (data segment, relation info not needed in executable.)
basinorminal	.line .strtab .Section header table	type defs global vars defined. (Syntab & debugging info are not loaded into memory)
of basi		Mapping between line no. in .c file and .Section present when -g.
		Sequence of String & NULL terminated String for Syntab, debug, section names in section header.

1. ELF Header :

- It contains metadata about the file, such as the architecture, the entry point address, the program header table location, the section header table location.
- It helps the loader understand how to load and execute the file.

2. Section Headers :

These describe various sections of the ELF file.

• **.text**: This section contains the executable code (program instructions) of the program.

• **.rodata**: Read-Only data, such as constants, strings or other constants that cannot be modified during execution.

• **.data**: This section contains initialized global and static variables that are read/write.

(program ofai bbaol (917 adt m)

• **.bss**: This section contains uninitialized global and static variables.

These variables are uninitialized to zero when the program starts.

• **Symtab**: It stands for Symbol table,

It is sometimes used for debugging and contains symbol information for the compiler or debugger.

• **.reltext**: It contains relocation entries for the .text section.

Used to adjust code references during linking.

• Stores relocation info not needed in executable.

• **.reldata**: It contains relocation entries for the data section.

- **.debug**: Contains debugging information, such as line numbers, variable names and other metadata helpful for debugging like type defs.
- **.line**: This section is typically used to store line number information used by debuggers to relate machine code back to source code.
- **.Strtab**: The String table, which contains all string literals used in the file like function names, variable names etc.
- **.T**: Stores null terminating strings used by various sections such as .symbol, .debug etc. Section names do not have a header.

3. Section Headers Table

This is a table of sections that describes each section's characteristics, such as its type, size and where it is located in the file.

- Each section header describes a specific section.

- * **Types of ELF (Object Files)**
 1. Relocatable Object Files
 - File extension is .o or .obj
 - They are produced by compilers and

are used to create executables or shared libraries after being linked.

- They contain machine code, but cannot be executed directly because they are incomplete. (They are intermediate output)

of compilation process.

• have Section headers.

2. Executable(ELF file)

- File extension usually no extension and often just file name like a.out or no custom extension like .elf .exe etc.

- They are the final outputs produced after linking all object files.

• They have program headers.

3. Shared Object file : (SO file)

- File extension is .so .dll .so .obj

- They are dynamically loaded libraries that can be shared amongst multiple processes.

- It has both Section headers and program headers.

- It contains code and data that can be shared by different programs.

4. Core dump file

- When program crashes the data is stored in core dump file where we can check what went wrong.

- Primarily used for debugging.